PROJECT REPORT: LAPTOP PRICE PREDICTION ENGINE

Subject: Data Science & Predictive Analytics

Name: Abhinav Jain

Registration number: 25BAI10234

Date: March 31, 2026

Course Name : Fundamentals in AI&ML

1. Executive Summary

This project focuses on the development of a predictive model that estimates the market price of a laptop based on its hardware specifications. By analyzing a dataset of over 1,300 laptop configurations, the system identifies how features like RAM, CPU clock speed, GPU brand, and screen resolution impact the final retail price. The goal was to move beyond "best guesses" and provide a data-driven tool for consumers and retailers.

---

2. Research & System Architecture

2.1 The Prediction Pipeline

The project follows a standard **Supervised Learning** workflow:

1. **Data Extraction:** Loading the raw CSV containing brand, type, RAM, weight, and specs.
2. **Feature Engineering:** Converting raw text (e.g., "8GB") into numerical values (e.g., 8).
3. **One-Hot Encoding:** Converting categorical data (Brand names like Apple, HP, Dell) into binary vectors so the math models can process them.
4. **Log Transformation:** Applying np.log to the price to normalize the distribution, as laptop prices are often "skewed" (many cheap ones, few very expensive ones).

2.2 Algorithm Selection: Random Forest Regressor

I chose the **Random Forest** algorithm because it is an "Ensemble" method. It creates hundreds of different "Decision Trees" and averages their results. This prevents "Overfitting," which is when a model memorizes the training data but fails to predict prices for new, unseen laptops.

---

3. Mathematical Methodology

3.1 Linear Relationship & Correlation

The model calculates the correlation between independent variables (X) and the dependent variable (Y, Price).

3.2 Evaluation Metric: R-Squared (R^2) and MAE

To judge the accuracy, I used:

- **MAE (Mean Absolute Error):** Telling us, on average, how many dollars/rupees the prediction was off by.
- **R^2 Score:** Measuring how much of the price variance is explained by the model (Targeting > 0.80).

---

4. Deep-Dive: Difficulties Faced & Technical Solutions

Predicting prices was significantly more difficult than simple classification due to these bottlenecks:

⚠️ Difficulty 1: Complex String Splitting (The "CPU" Problem)

Problem: The CPU column contained strings like "Intel Core i7 7700HQ 2.8GHz." The model cannot do math on a sentence.

Solution: I used **Regular Expressions (Regex)** to extract the CPU name and the Clock Speed separately. I then categorized CPUs into "Intel Core i7," "Other Intel," and "AMD" to simplify the model's decision-making process.

⚠️ Difficulty 2: Screen Resolution & Touchscreen Logic

Problem: Resolution was hidden inside text like "IPS Panel Retina Display 2560x1600."

Solution: I created two new binary features: Touchscreen (1 or 0) and IPS (1 or 0). Then, I extracted the X and Y resolutions to calculate **PPI (Pixels Per Inch)**, which is a much stronger predictor of price than raw resolution alone.

⚠️ Difficulty 3: Handling Categorical "Outliers"

Problem: Brands like "Razer" or "Apple" have much higher prices than "Acer" for the same specs because of brand value. A standard linear model failed to capture this.

Solution: I used a **Pipeline with One-Hot Encoding**. This allowed the model to treat each brand as a unique weight, essentially teaching the AI that "Brand Logo" is a feature just as important as "RAM."

⚠️ Difficulty 4: Price Skewness

Problem: The price data had a long "tail" (a few laptops cost $5,000, while most cost $500). This confused the model's accuracy.

Solution: I applied a **Logarithmic Transformation** to the target variable. This compressed the price range, making it easier for the algorithm to learn patterns. I then used np.exp to convert the prediction back to a normal price for the user.

---

5. Visual Analysis (EDA)

The project included several key visualizations:

- **Correlation Heatmap:** Proving that RAM and CPU have the highest impact on price.
- **Price vs Brand Bar Chart:** Showing the luxury premium of certain brands.
- **Scatter Plots:** Visualizing how the weight of a laptop inversely correlates with price (Thinner/lighter laptops are usually more expensive).

---

6. Future Scope & Conclusion

This project proves that hardware specifications can accurately predict market value within a 10% error margin. Future versions could include "Current Market Demand" or "Release Year" to make the predictions even more accurate.

Final Conclusion: By combining rigorous data cleaning with the Random Forest algorithm, I successfully built a tool that can assist buyers in determining if a laptop is "Overpriced" or a "Good Deal."